

# SUPERVISED AND EXPERIENTIAL LEARNING



## (Part VI – Ensemble of Classifiers/Multiple Classifiers /Diversity)

**Miquel Sànchez-Marrè**

Intelligent Data Science and Artificial Intelligence Research Centre (IDEAI-UPC)

Knowledge Engineering and Machine Learning Group (KEMLG-UPC)

Computer Science Dept.


Universitat Politècnica de Catalunya · BarcelonaTech

[miquel@cs.upc.edu](mailto:miquel@cs.upc.edu)

<http://www.cs.upc.edu/~miquel>

Course 2019/2020

<https://kemlg.upc.edu>




# PART 6 – ENSAMBLE OF CLASSIFIERS, MULTIPLE CLASSIFIERS, DIVERSITY

<https://kemlg.upc.edu>



Knowledge Engineering and Machine Learning Group  
UNIVERSITAT POLITÈCNICA DE CATALUNYA




# CLASSIFIER ENSEMBLE METHODS

<https://kemlg.upc.edu>



Knowledge Engineering and Machine Learning Group  
UNIVERSITAT POLITÈCNICA DE CATALUNYA





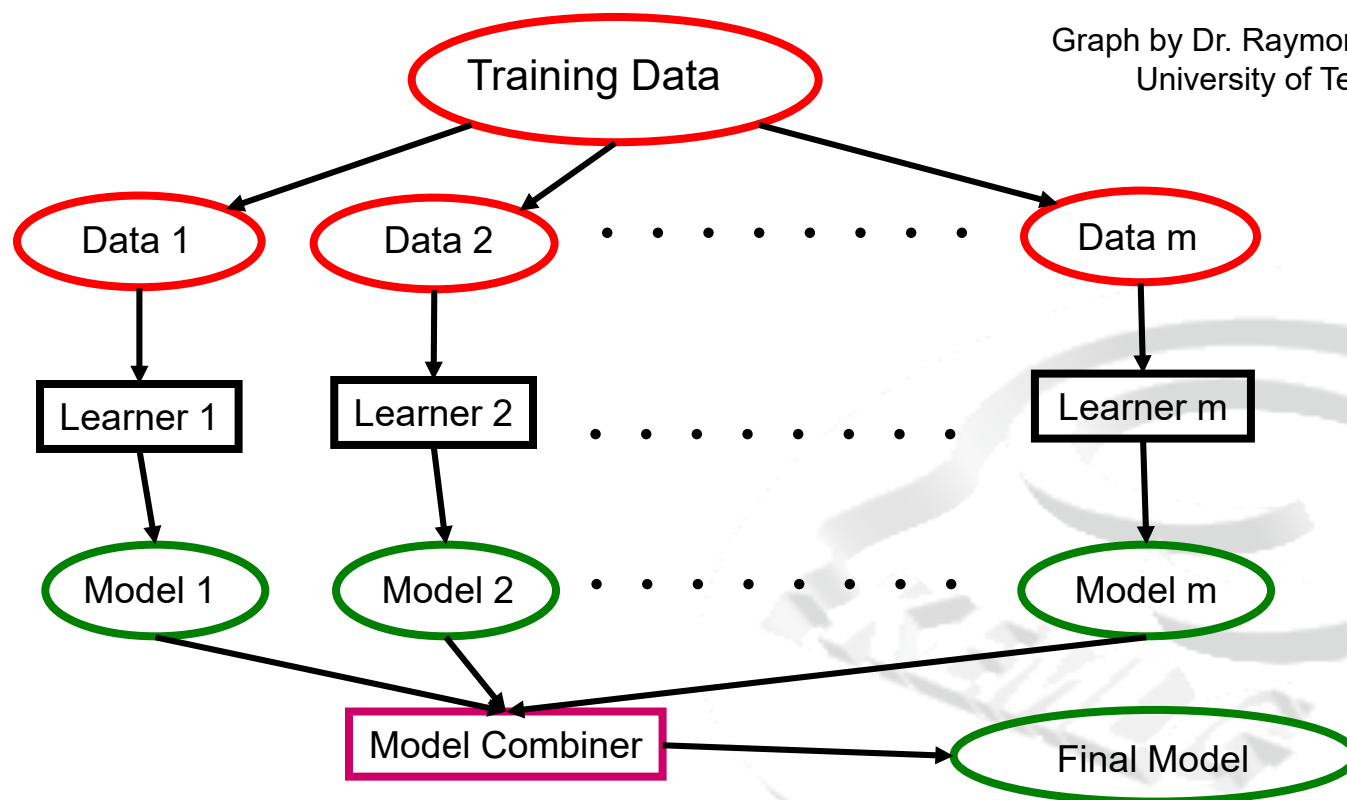
# Classifier Ensemble Methods

- **Goal:** to induce a *set of discriminant models (classifiers)*, which can be of the same type or not, from different samples or weighting the samples or randomly selecting the attributes at each split node of a tree, of a supervised database, aiming at *reducing the discrimination error* of each one of the (weak) classifiers .
- **Applicability criteria:** a *supervised database*, with one qualitative attribute being the *class attribute*, with a representative number of examples of the different possible *class labels*.
- **Most common methods:**
  - Bagging [Breiman *et al.*, 1984]
  - Boosting [Shapire, 1990], AdaBoost [Freund & Shapire, 1996]
  - Random Decision Forests [Ho, 1995], Decision Forests [Ho, 1998], Random Forests [Breiman, 2001]
- **Input:** original supervised data matrix
- **Output:** a *set of classifiers* which are able to discriminate/classify the qualitative attribute of interest (class attribute)
- **Evaluation Parameters:** predictive accuracy, scalability, robustness
- **Discrimination/classification process:** when a new instance must be discriminated, *all the classifiers* are used to get their class label prediction. The *class label assigned* to the new instance is *the result of a (weighted) voting among the class labels of the set of classifiers*



# Learning Ensembles

- Learn a set of discriminant/classifiers models using *different training data* or/and *different learning algorithms* and/or using other *diversification techniques*
- Combine the output of the classifiers, *i.e.* predicted label, using (weighted) majority voting scheme



Graph by Dr. Raymond J. Mooney  
University of Texas at Austin



# Value of Ensembles


- When combining multiple *independent* and *diverse* decisions:
  - Each decision is more accurate than random guessing
  - Random errors cancel each other out
  - Correct decisions are reinforced
  - Classification accuracy increases





# Different types of ensemble learning

- Different learning **algorithms**
- Algorithms with different choice for **hyper-parameters**
- Data set with different **features** (e.g. random subspace)
- Data set = different **subsets** (e.g. bagging, boosting)







# A possible classification

- Multiexpert combination methods (parallel classifier models):
  - Global approach (classifier models' fusion)
    - ◆ **Voting**: voting among different classifiers
    - ◆ **Bagging**: resample training data (bootstrapping), same classifier and voting
    - ◆ **Stacking**: a combiner learner (meta-learner), which combines the predictions of the classifiers
    - ◆ **Randomizing input features**: random subsets of features at each node (random forests)
  - Local approach (classifier models' selection)
    - ◆ **Gating**: Meta selection of the best classifier/s (best local expert/s) to be used (*Mixture of experts ensemble*)
- Multistage combination (sequential classifier models)
  - **Boosting**: Reweight training data. Next Learner focusing on misclassified instances by previous classifier
  - **Cascading**: Increasing complexity of learners
- Decorating methods: Adding artificial training data (noise addition)



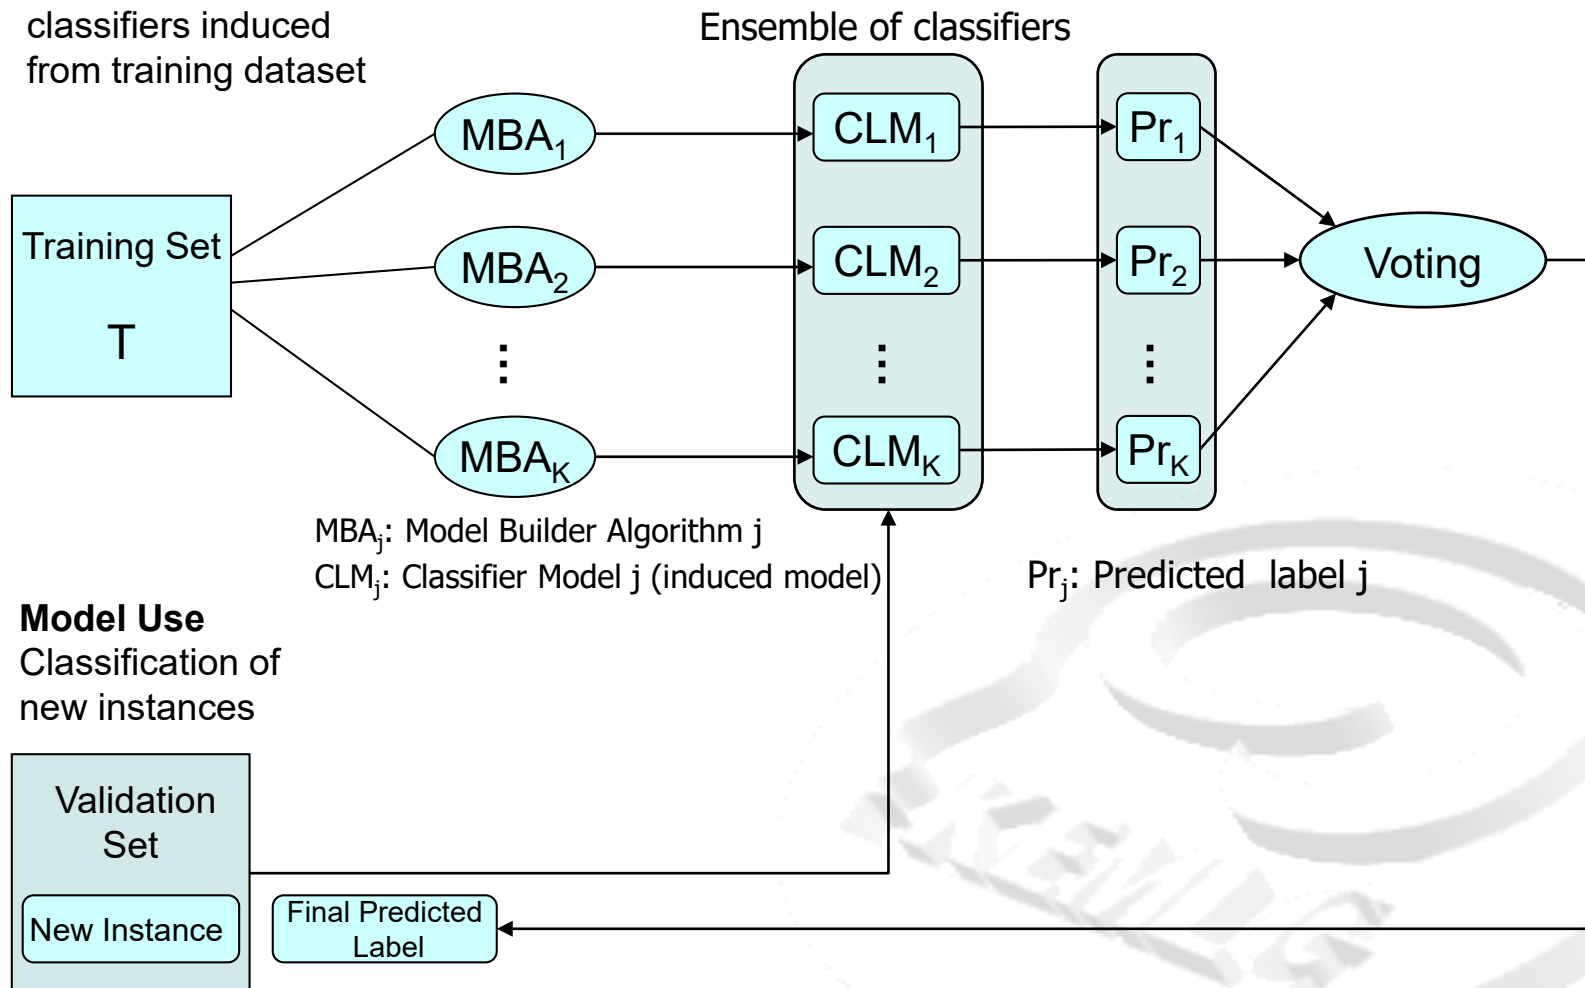


# Voting

- Different algorithms, same set of training data

## Model Building

Ensemble of classifiers induced from training dataset





# Bagging (Bootstrap Aggregating) (1)

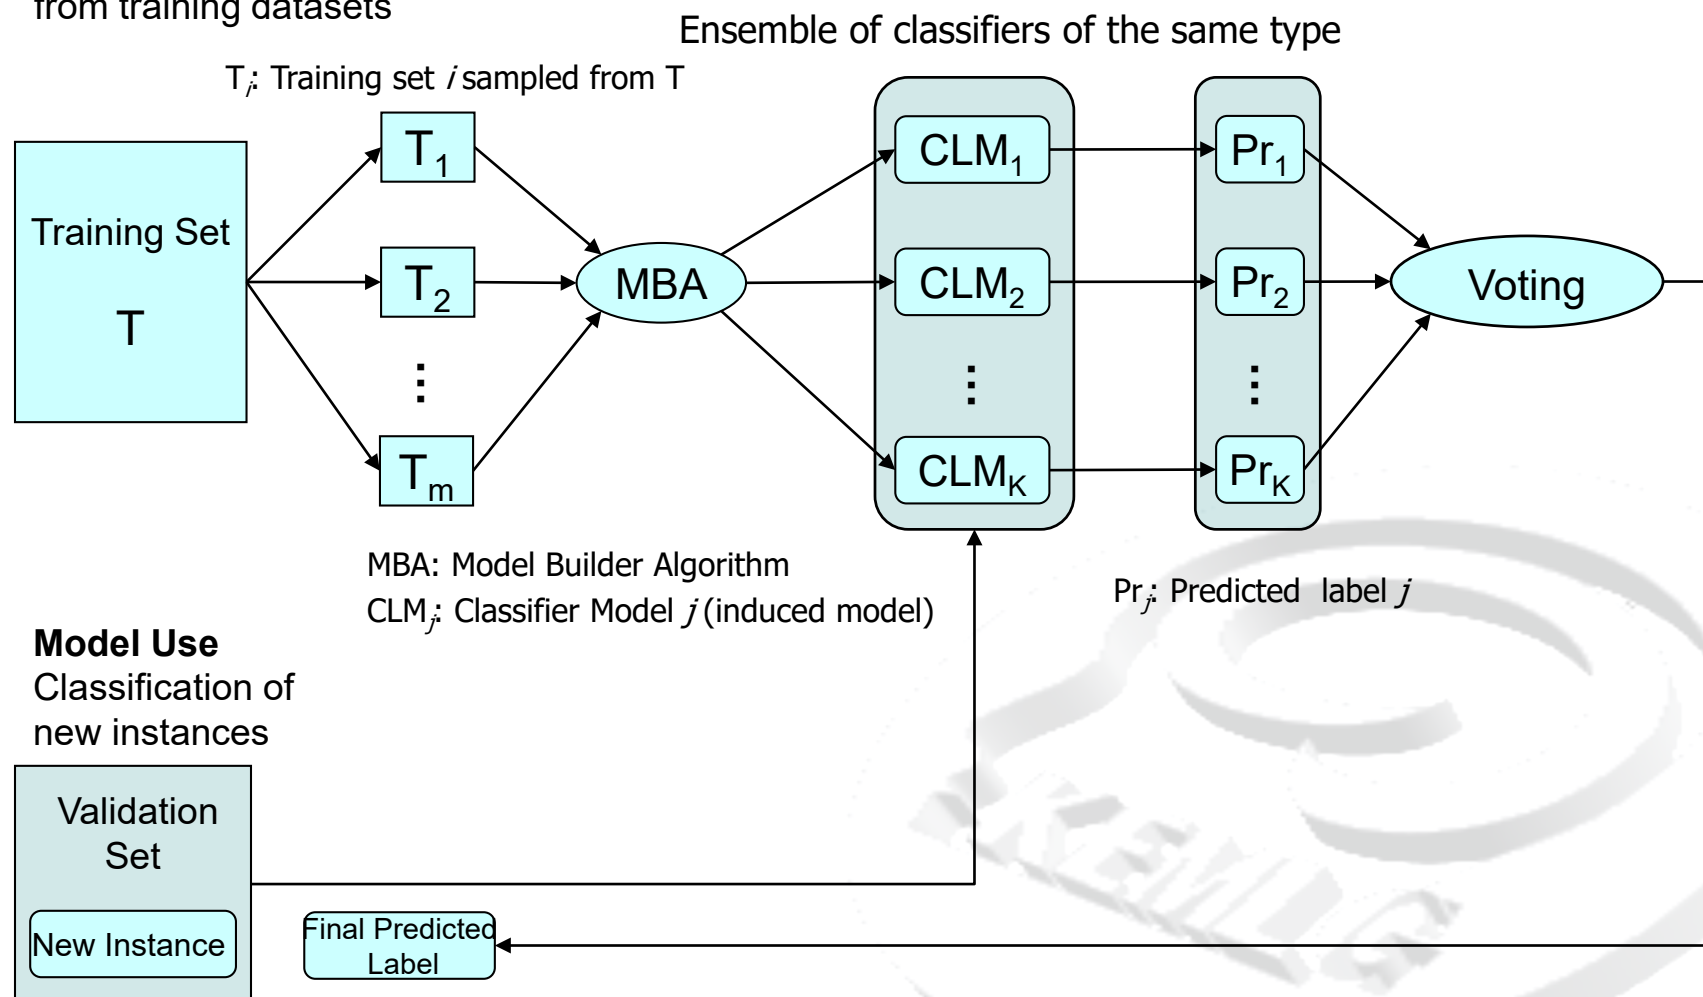
- Bootstrap aggregating. Create ensembles by repeatedly randomly resampling the training data [Breiman, 1996].
- Given a training set of size  $n$ , create  $m$  samples of size  $n$  by drawing  $n$  examples from the original data, *with replacement*.
  - Each *bootstrap sample* will on average contain 63.2% of the unique training examples, the rest are replicates.
- Combine the  $m$  resulting models using simple majority vote.
- Decreases error by decreasing the variance in the results due to *unstable learners*, algorithms (like decision trees) whose output can change dramatically when the training data is slightly changed.

# Bagging (Bootstrap Aggregating) (2)

- Same algorithm, different versions of training dataset:

## Model Building

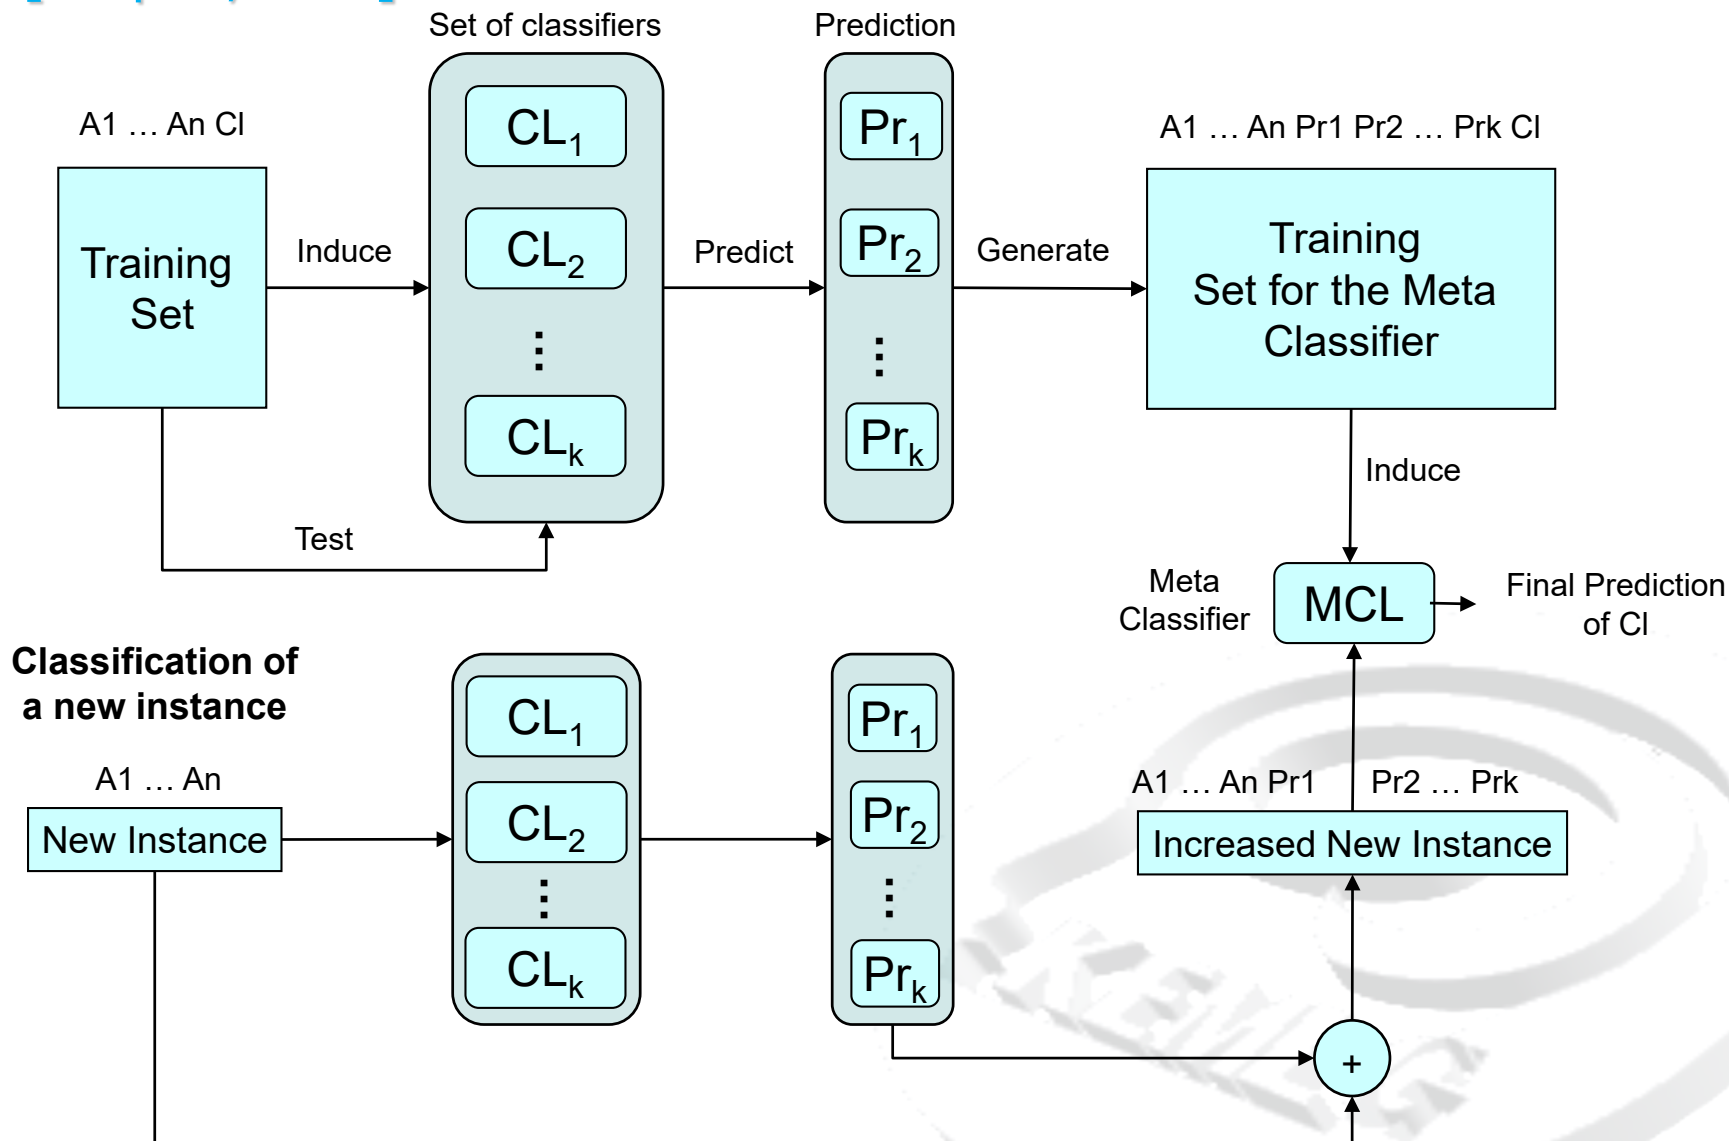
Ensemble of classifiers induced from training datasets





# Stacking (Stacked Generalization)

[Wolpert, 1992]



# Random Decision Forests / Random Forests

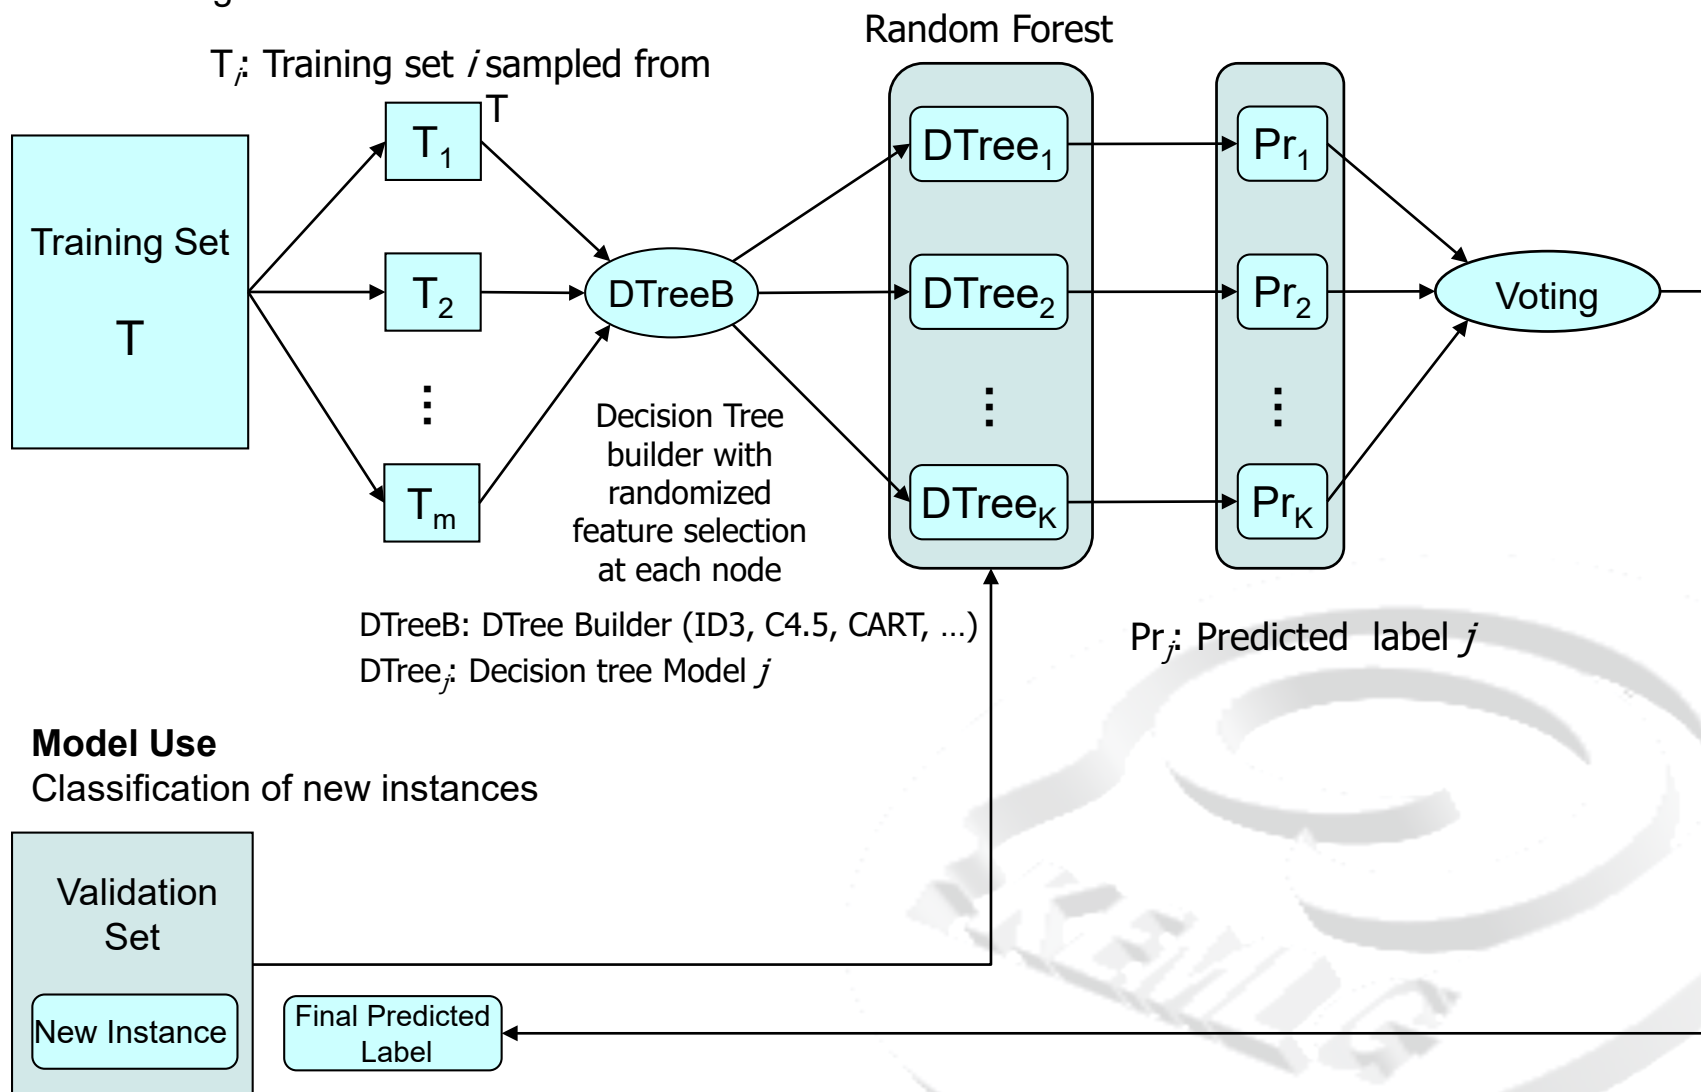
[Ho, 1998] / [Breiman, 2001]

- References:
  - Tin Kam Ho. The Random Subspace Method for Constructing Decision Forests. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 20(8):832-844, 1998.
  - Leo Breiman. Random Forests. *Machine Learning*, 45:5-32, 2001
- Motivation: reduce error correlation between classifiers
- Main idea: build a larger number of un-pruned decision trees
- Ho's proposal: each tree is grown using a random subspace (selection) of features, which is the same for all the node splits
- Breiman's proposal: each tree uses a random subspace (selection) of features to split on at each node, and the training set for each tree is sampled (bootstrapping) from the original dataset



## Model Building

Ensemble of classifiers induced from training datasets



# How Random Forests Work

[Breiman, 2001]

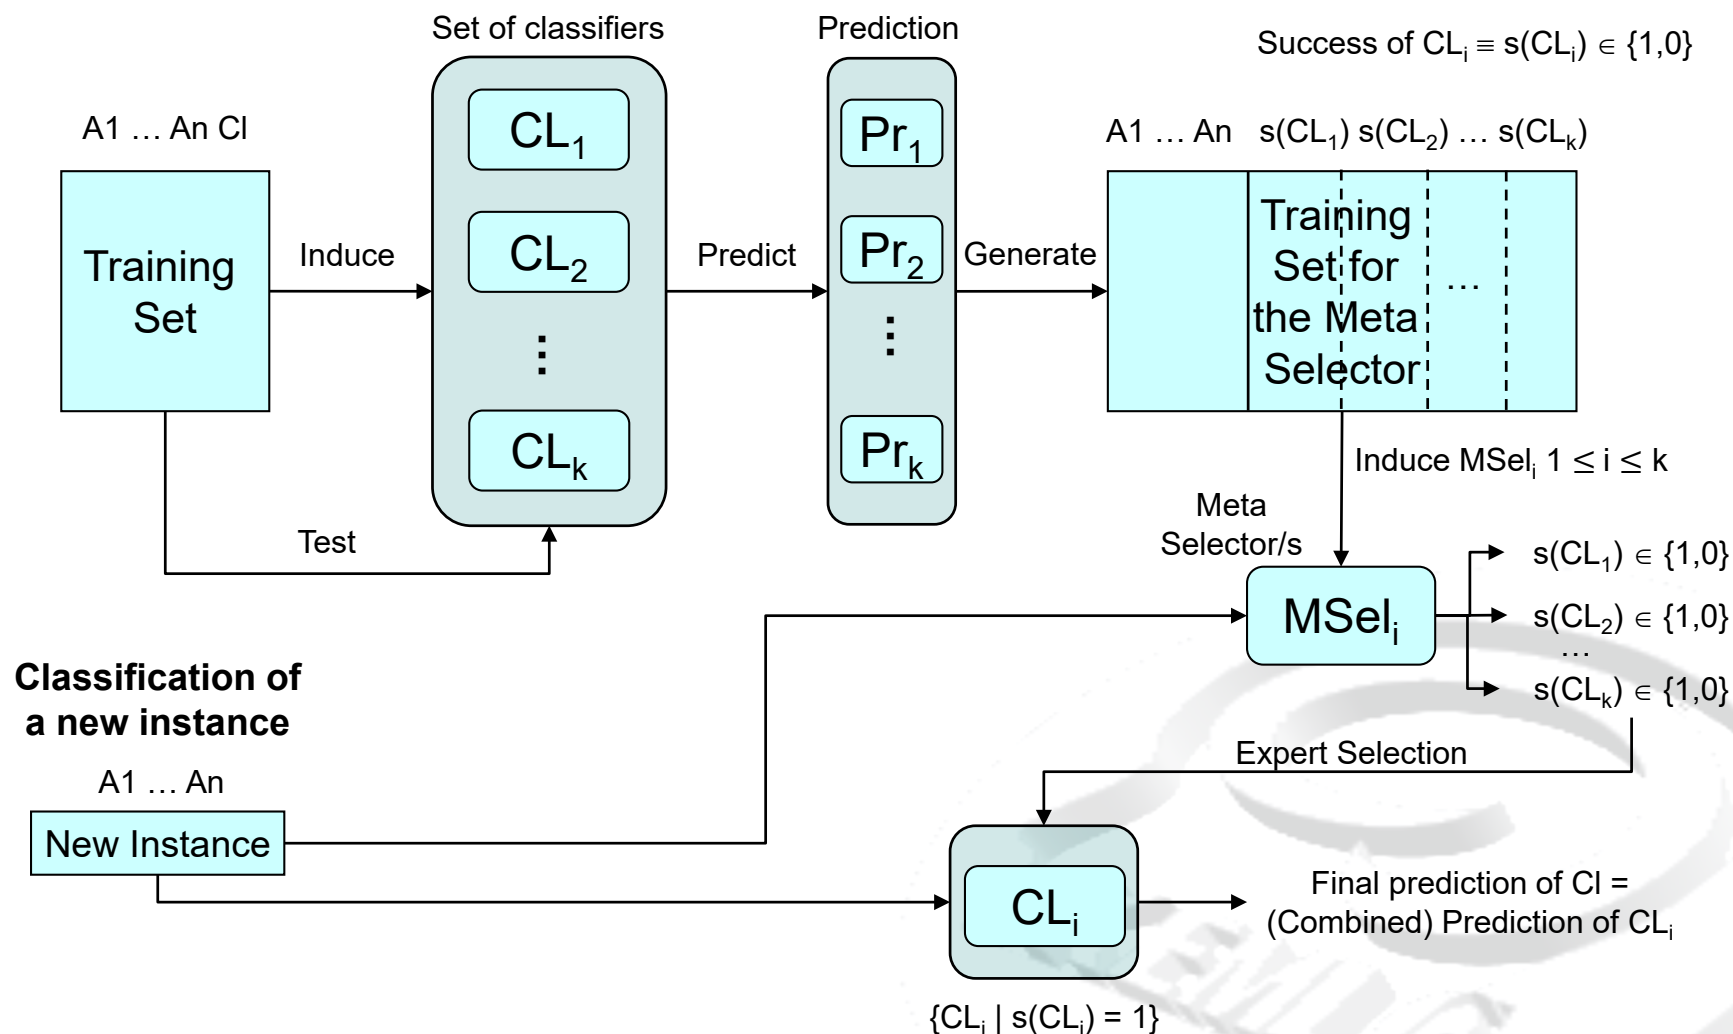
- Each tree is grown on a bootstrap sample of the training set of **N** cases.
- A number **F** is specified much smaller than the total number of variables **M**:
  - $F = \text{sqrt}(M)$  or
  - $F = \text{int}(\log_2 M + 1)$
- At each node,  $F$  variables are selected at random out of the  $M$ .
- The split used is the best split on these **F** variables according to the decision tree strategy.
- Final classification is done by majority vote across trees.





# Gating (Mixture of Experts/Experts' Selection)

[Jacobs et al., 1991]





# Boosting (1)

[Schapire, 1990]

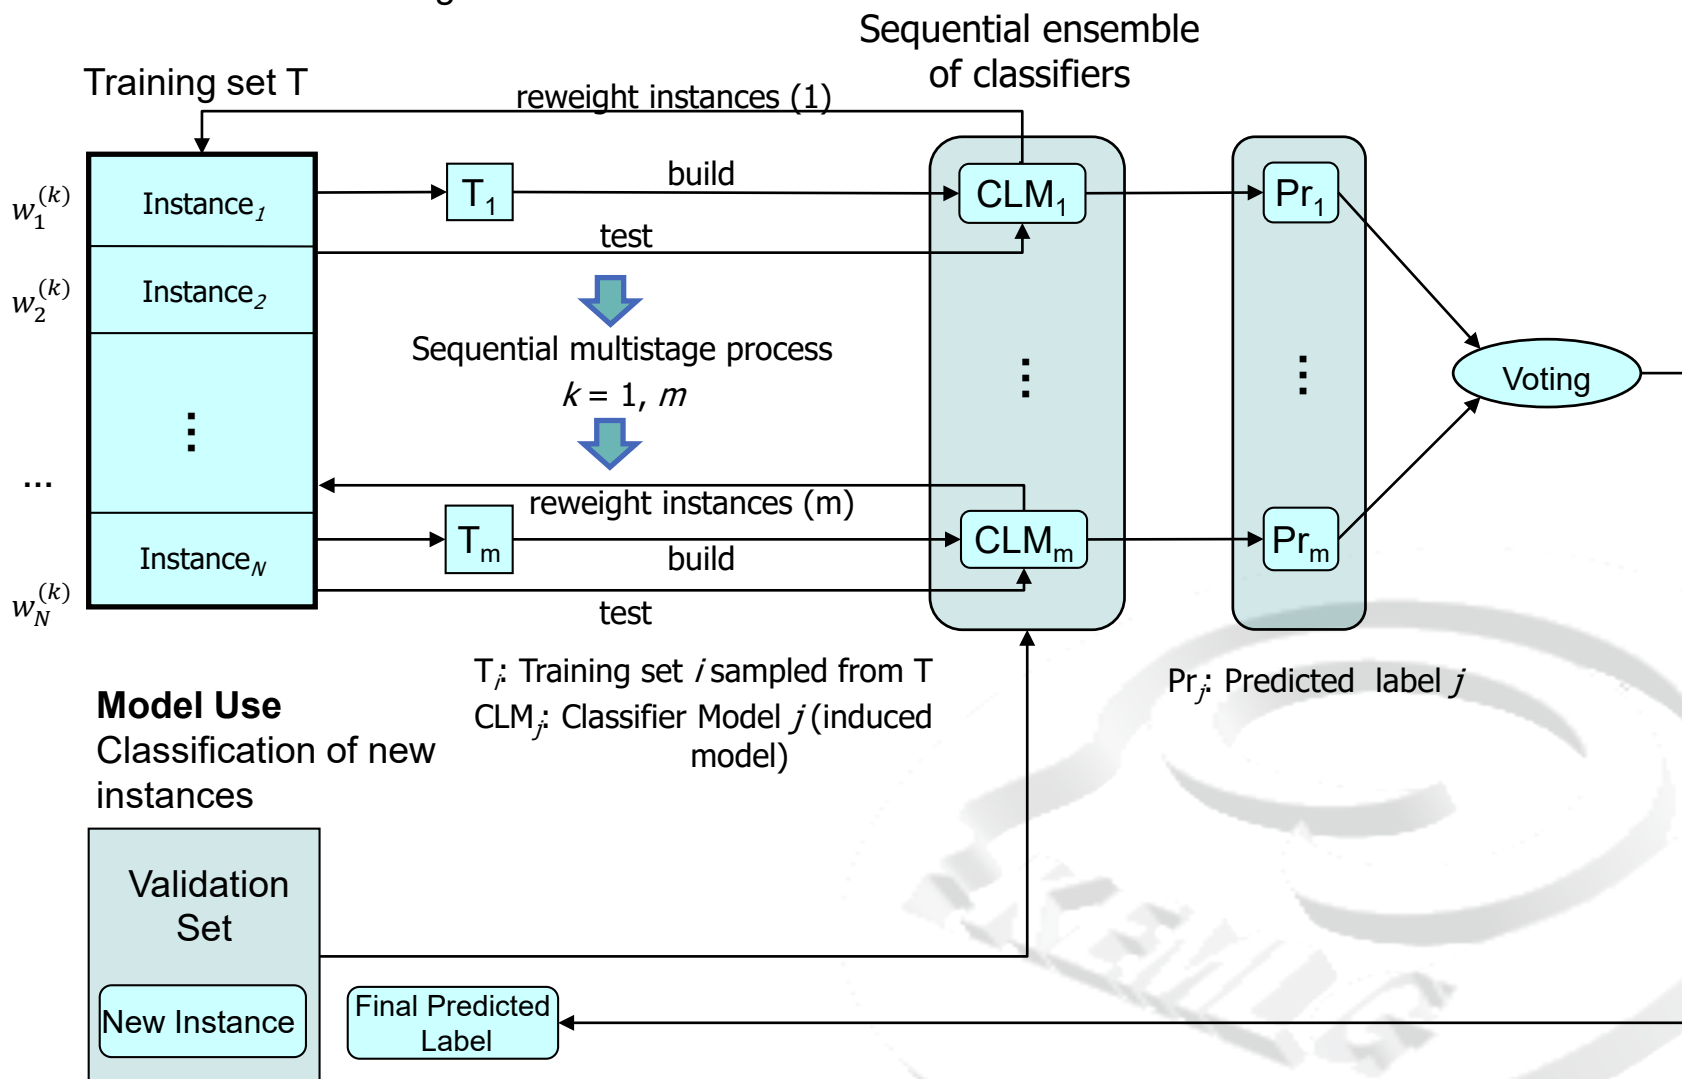
- Originally developed by computational learning theorists to guarantee *performance improvements* on fitting training data for a **weak learner** that only needs to generate a hypothesis with a training accuracy greater than 0.5 [Schapire, 1990].
- Revised to be a practical algorithm, **AdaBoost**, for building ensembles that empirically improves generalization performance [Freund & Shapire, 1996].
- Examples are given weights. At each iteration, a new hypothesis is learned and the examples are reweighted to focus the system on examples that the most recently learned classifier got wrong.



# Boosting (2)

## Model Building

Ensemble of classifiers sequentially induced from training datasets





# Boosting: basic algorithm

- General boosting algorithm:

Set all examples to have equal uniform weights

**for**  $t$  from 1 to  $T$  **do**

Learn a hypothesis/model,  $h_t$  from the weighted examples

Decrease the weights of examples  $h_t$  classifies correctly

**endfor**

- Base (weak) learner must **focus on** *correctly classifying the most highly weighted examples* while strongly avoiding over-fitting.
- During *testing*, each of the  $T$  hypotheses/models get a weighted vote proportional to their accuracy on the training data.

# AdaBoost

[Freund & Shapire, 1996]

TrainAdaBoost( $D$ , BaseLearn)

**for each** example  $d_i$  in  $D$  **do**

    let its weight  $w_i = 1/|D|$

**endforeach**

Let  $H$  be an empty set of hypotheses

**for**  $t$  from 1 to  $T$  **do**

    Learn a hypothesis,  $h_t$ , from the weighted examples:  $h_t = \text{BaseLearn}(D)$

    Add  $h_t$  to  $H$

    Calculate the error,  $\epsilon_t$ , of the hypothesis  $h_t$  as the total sum weight of the examples that it classifies incorrectly

**If**  $\epsilon_t > 0.5$  **then** exit loop, **else** continue

    Let  $\beta_t = \epsilon_t / (1 - \epsilon_t)$

    Multiply the weights of the examples that  $h_t$  classifies correctly by  $\beta_t$

    Rescale the weights of all of the examples so the total sum weight remains 1.

**endfor**

**return**  $H$

TestAdaBoost( $ex$ ,  $H$ )

    Let each hypothesis,  $h_t$ , in  $H$  vote for  $ex$ 's classification with weight  $\log(1/\beta_t)$

**return** the class with the highest weighted vote in total.



# Learning with Weighted Examples

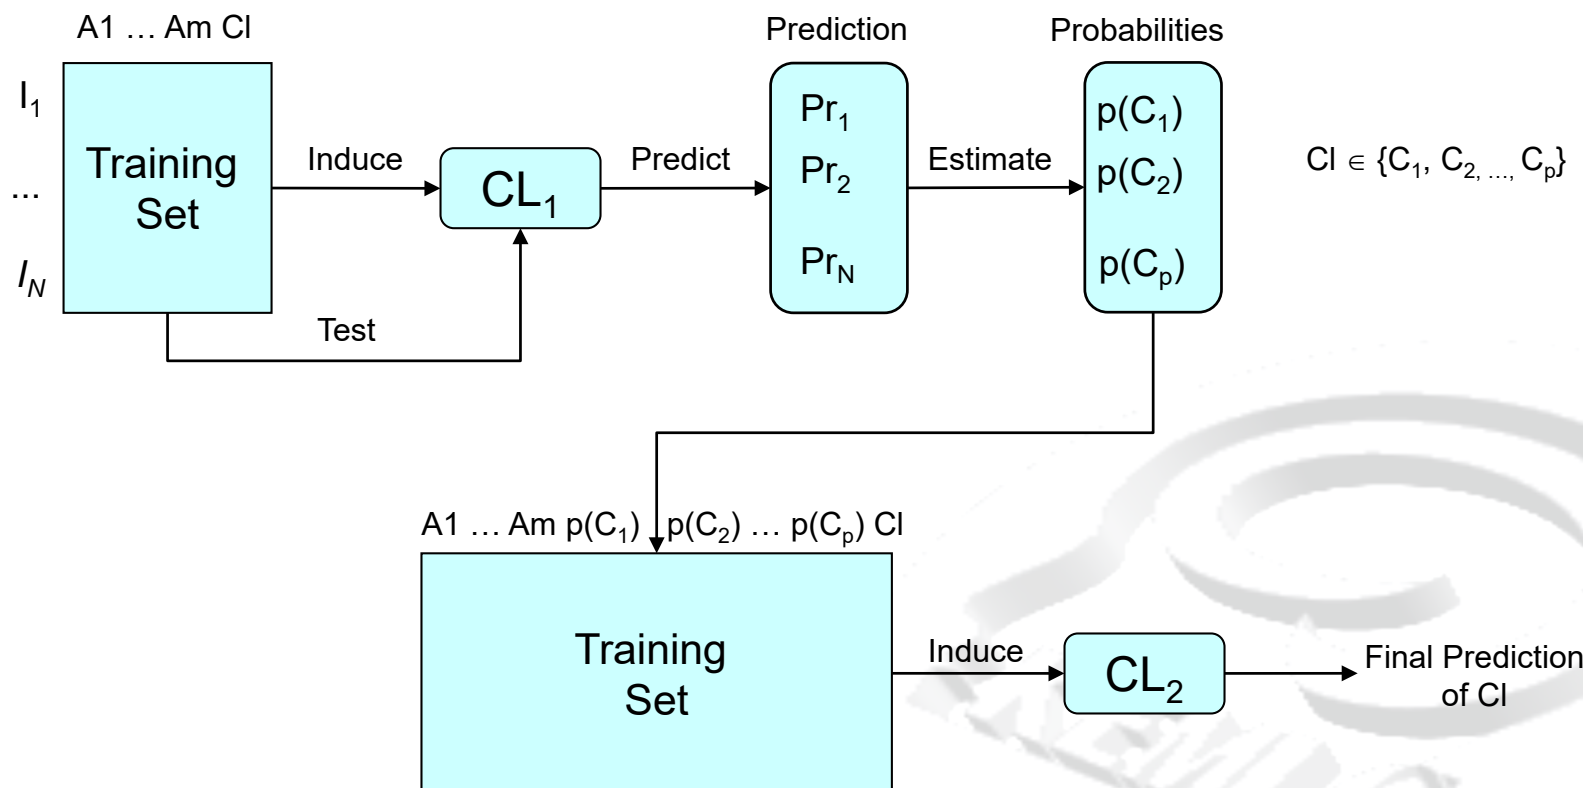
- Generic approach is to *replicate examples in the training set proportional to their weights*
  - $\frac{w_i}{\sum w_i} = \frac{n_i}{N}$  , thus  $n_i = w_i * N$
  - e.g., if we have a total number of  $N=1000$  examples, in the weighted sample there should be:
    - ◆ 10 replicates of an example with a weight of 0.01
    - ◆ 100 replicates of one example with weight 0.1
- Most algorithms can be enhanced to efficiently incorporate weights directly in the learning algorithm so that the effect is the same (e.g. implement the WeightedInstancesHandler interface in WEKA).
- For decision trees, for calculating information gain, when counting example  $i$ , simply increment the corresponding count by  $w_i$  rather than by 1.



# Cascading

[Viola & Jones, 2001]

- Sequence of several classifiers, using all information collected from the output from a previous classifier as additional information for the next classifier in the *cascade*







# Experimental Results on Ensembles

[Freund & Schapire, 1996; Quinlan, 1996]

- Ensembles have been used to *improve generalization accuracy* on a wide variety of problems.
- On average, *Boosting* provides a larger increase in accuracy than *Bagging*.
- *Boosting* on rare occasions can degrade accuracy.
- *Bagging* more consistently provides a modest improvement.
- *Boosting* is particularly subject to over-fitting when there is significant noise in the training data.
- *Bagging* is easily parallelized.
- *Boosting* is not easily parallelized.



# Random Forests vs Adaboost

- Error rates compare favorably to Adaboost
- More robust with respect to noise.
- More efficient on large data
- Provides an estimation of the importance of features in determining classification



# Ensemble Methods

- RapidMiner operators:
  - Modeling/Predictive/Ensembles:
    - ◆ Bagging
    - ◆ Adaboost
    - ◆ Vote (different classifiers)
    - ◆ Stacking (different classifiers training a high level classifier)
    - ◆ ...
  - Modeling/Predictive/Trees:
    - ◆ Random Forest
- Python
  - Scikit Learning
    - ◆ BaggingClassifier
    - ◆ RandomForestClassifier
    - ◆ ....
- R
  - Caret package





# Intelligent Data Science and Artificial Intelligence (IDEAI-UPC)

**Miquel Sànchez-Marrè**  
**miquel@cs.upc.edu**



Knowledge Engineering and Machine Learning Group  
UNIVERSITAT POLITÈCNICA DE CATALUNYA

<https://kemlg.upc.edu>